

분산시스템 과제 #6

박종현, 2025-05-09

공과대학 컴퓨터정보통신공학과

과제 목표

1. 3-5 자료 : 4,5,6,7,9,10,12,13쪽
2. 3-6 자료 : 7~13쪽
3. 3-7 자료 : 3~7쪽
4. 3-8 자료 : 3~10쪽

[ch 3-9]

```
INFO : The url to track the job: http://DESKTOP-635D6IV.:50000/proxy/application_1746737381207_0032/
INFO : Starting Job = job_1746737381207_0032, Tracking URL = http://DESKTOP-635D6IV.:50000/proxy/application_1746737381207_0032/
INFO : Kill Command = /home/hduser/hadoop_test/hadoop-3.4.0/bin/mapred job -kill job_1746737381207_0032
INFO : Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
INFO : 2025-05-09 06:23:42,710 Stage-1 map = 0%, reduce = 0%
INFO : 2025-05-09 06:23:48,860 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 3.46 sec
INFO : 2025-05-09 06:23:55,039 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 7.18 sec
INFO : MapReduce Total cumulative CPU time: 7 seconds 180 msec
INFO : Ended Job = job_1746737381207_0032
INFO : MapReduce Jobs Launched:
INFO : Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 7.18 sec HDFS Read: 16659 HDFS Write: 172 HDFS EC Read: 0 S
SUCCESS
INFO : Total MapReduce CPU Time Spent: 7 seconds 180 msec
INFO : Completed executing command(queryId=hduser_20250509062333_916d3f1c-48a9-4eb3-a6a6-091cea07208c); Time taken: 22.378 seconds

+-----+
| avg_salary |
+-----+
| 6500       |
| 6500       |
| 6500       |
| 6500       |
| 6500       |
+-----+
5 rows selected (22.565 seconds)
0: jdbc:hive2://localhost:10000>
```

```
INFO : set mapreduce.job.reduces=<number>
INFO : number of splits:1
INFO : Submitting tokens for job: job_1746737381207_0033
INFO : Executing with tokens: []
INFO : The url to track the job: http://DESKTOP-635D6IV.:50000/proxy/application_1746737381207_0033/
INFO : Starting Job = job_1746737381207_0033, Tracking URL = http://DESKTOP-635D6IV.:50000/proxy/application_1746737381207_0033/
INFO : Kill Command = /home/hduser/hadoop_test/hadoop-3.4.0/bin/mapred job -kill job_1746737381207_0033
INFO : Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
INFO : 2025-05-09 06:24:21,857 Stage-1 map = 0%, reduce = 0%
INFO : 2025-05-09 06:24:26,993 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 3.02 sec
INFO : 2025-05-09 06:24:33,157 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 7.42 sec
INFO : MapReduce Total cumulative CPU time: 7 seconds 420 msec
INFO : Ended Job = job_1746737381207_0033
INFO : MapReduce Jobs Launched:
INFO : Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 7.42 sec HDFS Read: 18362 HDFS Write: 115 HDFS EC Read: 0 S
SUCCESS
INFO : Total MapReduce CPU Time Spent: 7 seconds 420 msec
INFO : Completed executing command(queryId=hduser_20250509062410_6ae7910e-f848-4458-bc0b-5ba48fe657c3); Time taken: 24.4 seconds

+-----+ +-----+ +-----+
| avg_window_0 | min_window_1 | max_window_2 |
+-----+ +-----+ +-----+
| 1300.0       | 800          | 2000         |
+-----+ +-----+ +-----+
1 row selected (24.534 seconds)
0: jdbc:hive2://localhost:10000>
```

```

207_0035/
INFO : Kill Command = /home/hduser/hadoop_test/hadoop-3.4.0/bin/mapred job -kill job_1746737381207_0035
INFO : Hadoop job information for Stage-2: number of mappers: 1; number of reducers: 1
INFO : 2025-05-09 06:25:27,486 Stage-2 map = 0%, reduce = 0%
INFO : 2025-05-09 06:25:33,731 Stage-2 map = 100%, reduce = 0%, Cumulative CPU 3.4 sec
INFO : 2025-05-09 06:25:40,969 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 9.09 sec
INFO : MapReduce Total cumulative CPU time: 9 seconds 90 msec
INFO : Ended Job = job_1746737381207_0035
INFO : MapReduce Jobs Launched:
INFO : Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 6.27 sec HDFS Read: 16374 HDFS Write: 338 HDFS EC Read: 0 S
SUCCESS
INFO : Stage-Stage-2: Map: 1 Reduce: 1 Cumulative CPU: 9.09 sec HDFS Read: 15207 HDFS Write: 310 HDFS EC Read: 0 S
SUCCESS
INFO : Total MapReduce CPU Time Spent: 15 seconds 360 msec
INFO : Completed executing command(queryId=hduser_20250509062450_87ae73dc-4991-4581-9208-506478fb1a77); Time taken: 51.002 seconds

+----+-----+-----+-----+-----+-----+
| name | salary | projects | avg_salary | rank | prank |
+----+-----+-----+-----+-----+-----+
| paul | 2000 | ["proj1","proj2"] | 1300.0 | 1 | 0.0 |
| david | 1500 | ["proj1"] | 1300.0 | 2 | 0.25 |
| nina | 1200 | ["proj2"] | 1300.0 | 3 | 0.5 |
| peter | 1000 | ["proj1"] | 1300.0 | 4 | 0.75 |
| michael | 800 | ["proj1","proj2"] | 1300.0 | 5 | 1.0 |
+----+-----+-----+-----+-----+-----+

5 rows selected (51.212 seconds)
0: jdbc:hive2://localhost:10000>

```

```

INFO : Kill Command = /home/hduser/hadoop_test/hadoop-3.4.0/bin/mapred job -kill job_1746737381207_0037
INFO : Hadoop job information for Stage-2: number of mappers: 1; number of reducers: 1
INFO : 2025-05-09 06:26:37,025 Stage-2 map = 0%, reduce = 0%
INFO : 2025-05-09 06:26:43,218 Stage-2 map = 100%, reduce = 0%, Cumulative CPU 5.18 sec
INFO : 2025-05-09 06:26:50,422 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 8.82 sec
INFO : MapReduce Total cumulative CPU time: 8 seconds 820 msec
INFO : Ended Job = job_1746737381207_0037
INFO : MapReduce Jobs Launched:
INFO : Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 6.28 sec HDFS Read: 17036 HDFS Write: 338 HDFS EC Read: 0 S
SUCCESS
INFO : Stage-Stage-2: Map: 1 Reduce: 1 Cumulative CPU: 8.82 sec HDFS Read: 16033 HDFS Write: 338 HDFS EC Read: 0 S
SUCCESS
INFO : Total MapReduce CPU Time Spent: 15 seconds 100 msec
INFO : Completed executing command(queryId=hduser_20250509062600_6df50979-05e5-4ecd-bd6f-8d4c4a2772f0); Time taken: 51.079 seconds

+----+-----+-----+-----+-----+-----+
| name | salary | projects | avg_salary | rank | prank |
+----+-----+-----+-----+-----+-----+
| paul | 2000 | ["proj1","proj2"] | 1325.0 | 1 | 0.0 |
| david | 1500 | ["proj1"] | 1325.0 | 2 | 0.3333333333333333 |
| peter | 1000 | ["proj1"] | 1325.0 | 3 | 0.6666666666666666 |
| michael | 800 | ["proj1","proj2"] | 1325.0 | 4 | 1.0 |
| nina | 1200 | ["proj2"] | 1200.0 | 1 | 0.0 |
+----+-----+-----+-----+-----+-----+

5 rows selected (51.245 seconds)
0: jdbc:hive2://localhost:10000>
0: jdbc:hive2://localhost:10000>

```

```

+----+-----+-----+-----+-----+-----+
| isPivotResult: true |
+----+-----+-----+-----+-----+-----+
| Explain |
+----+-----+-----+-----+-----+-----+
| Statistics: Num rows: 1 Data size: 2108 Basic stats: COMPLETE Column stats: NONE |
| Select Operator |
| expressions: _col1 (type: string), _col2 (type: int), _col5 (type: array<string>), _col0 (type: double), |
| rank_window_1 (type: int), percent_rank_window_2 (type: double) |
| outputColumnNames: _col0, _col1, _col2, _col3, _col4, _col5 |
| Statistics: Num rows: 1 Data size: 2108 Basic stats: COMPLETE Column stats: NONE |
| File Output Operator |
| compressed: false |
| Statistics: Num rows: 1 Data size: 2108 Basic stats: COMPLETE Column stats: NONE |
| table: |
| input format: org.apache.hadoop.mapred.SequenceFileInputFormat |
| output format: org.apache.hadoop.hive.ql.io.HiveSequenceFileOutputFormat |
| serde: org.apache.hadoop.hive.serde2.lazy.LazySimpleSerDe |
|
| Stage: Stage-0 |
| Fetch Operator |
| limit: -1 |
| Processor Tree: |
| ListSink |
+----+-----+-----+-----+-----+-----+

119 rows selected (0.216 seconds)
0: jdbc:hive2://localhost:10000>

```

[ch 3-11]

```

hduser@ubuntu: ~
20% 14 GB 76%
hduser@ubuntu:~$ jps
1363 HMaster
1735 JarBootstrapMain
1947 Jps
hduser@ubuntu:~$ ss -ltn
State  Recv-Q  Send-Q      Local Address:Port      Peer Address:Port  Process
LISTEN  0        4096      127.0.0.53%lo:53        0.0.0.0:*
LISTEN  0        4096      127.0.0.54:53          0.0.0.0:*
LISTEN  0         128             *:16010                *:16010
LISTEN  0         128             *:16030                *:16030
LISTEN  0          50  [::ffff:127.0.0.1]:2181          :::*
LISTEN  0        4096  [::ffff:127.0.1.1]:16020          :::*
LISTEN  0        4096  [::ffff:127.0.1.1]:16000          :::*
LISTEN  0        4096             *:22                    *:22
hduser@ubuntu:~$

```

```

hduser@ubuntu: ~/hbase_test/hbase-2.6.2/bin
30% 13 GB 76%
hduser@ubuntu:~/hbase_test/hbase-2.6.2/bin$ ./hbase shell
HBase Shell
Use "help" to get list of supported commands.
Use "exit" to quit this interactive shell.
For Reference, please visit: http://hbase.apache.org/2.0/book.html#shell
Version 2.6.2, r6b3b36b429cf9a9d74110de79eb3b327b29ebf17, Fri Feb 14 14:39:16 UT
C 2025
Took 0.0011 seconds
hbase:001:0> create 'test', 'cf1'
Created table test
Took 0.9918 seconds
=> Hbase::Table - test
hbase:002:0> create 'test1', 'cf1', 'cf2'
Created table test1
Took 0.6605 seconds
=> Hbase::Table - test1
hbase:003:0> list
TABLE
test
test1
2 row(s)
Took 0.0246 seconds
=> ["test", "test1"]
hbase:004:0>

```

```

hduser@ubuntu: ~/hbase_test/hbase-2.6.2/bin
17% 14 GB 76%
hbase:004:0> describe 'test'
Table test is ENABLED
test, {TABLE_ATTRIBUTES => {METADATA => {'hbase.store.file-tracker.impl' => 'DEFAULT'}}}
COLUMN FAMILIES DESCRIPTION
{NAME => 'cf1', INDEX_BLOCK_ENCODING => 'NONE', VERSIONS => '1', KEEP_DELETED_CELLS => 'FALSE', DATA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER', MIN_VERSIONS => '0', REPLICATION_SCOPE => '0', BLOOMFILTER => 'ROW', IN_MEMORY => 'false', COMPRESSION => 'NONE', BLOCKCACHE => 'true', BLOCKSIZE => '65536 B (64KB)'}

1 row(s)
Quota is disabled
Took 0.1031 seconds
hbase:005:0> describe 'test1'
Table test1 is ENABLED
test1, {TABLE_ATTRIBUTES => {METADATA => {'hbase.store.file-tracker.impl' => 'DEFAULT'}}}
COLUMN FAMILIES DESCRIPTION
{NAME => 'cf1', INDEX_BLOCK_ENCODING => 'NONE', VERSIONS => '1', KEEP_DELETED_CELLS => 'FALSE', DATA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER', MIN_VERSIONS => '0', REPLICATION_SCOPE => '0', BLOOMFILTER => 'ROW', IN_MEMORY => 'false', COMPRESSION => 'NONE', BLOCKCACHE => 'true', BLOCKSIZE => '65536 B (64KB)'}

{NAME => 'cf2', INDEX_BLOCK_ENCODING => 'NONE', VERSIONS => '1', KEEP_DELETED_CELLS => 'FALSE', DATA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER', MIN_VERSIONS => '0', REPLICATION_SCOPE => '0', BLOOMFILTER => 'ROW', IN_MEMORY => 'false', COMPRESSION => 'NONE', BLOCKCACHE => 'true', BLOCKSIZE => '65536 B (64KB)'}

```

```

hduser@ubuntu: ~/hbase_test/hbase-2.6.2/bin
24% 14 GB 75%
'0', REPLICATION_SCOPE => '0', BLOOMFILTER => 'ROW', IN_MEMORY => 'false', COMPRESSION => 'NONE', BLOCKCACHE => 'true', BLOCKSIZE => '65536 B (64KB)'}

2 row(s)
Quota is disabled
Took 0.0331 seconds
hbase:006:0> put 'test', 'key1', 'cf1:a', 'value1'
Took 0.0603 seconds
hbase:007:0> put 'test', 'key2', 'cf1:a', 'value2'
Took 0.0097 seconds
hbase:008:0> put 'test', 'key2', 'cf1:a', 'value2'
Took 0.0049 seconds
hbase:009:0> put 'test', 'key4', 'cf1:c', 'value4'
Took 0.0076 seconds
hbase:010:0> scan 'test'
ROW COLUMN+CELL
key1 column=cf1:a, timestamp=2025-05-15T07:35:59.138, value=value1
key2 column=cf1:a, timestamp=2025-05-15T07:36:05.394, value=value2
key4 column=cf1:c, timestamp=2025-05-15T07:36:08.459, value=value4

3 row(s)
Took 0.0288 seconds
hbase:011:0>

```

```

hduser@ubuntu: ~/hbase_test/hbase-2.6.2/bin
22% 14 GB 75%
hbase:009:0> put 'test', 'key4', 'cf1:c', 'value4'
Took 0.0076 seconds
hbase:010:0> scan 'test'
ROW COLUMN+CELL
key1      column=cf1:a, timestamp=2025-05-15T07:35:59.138, value=val
          ue1
key2      column=cf1:a, timestamp=2025-05-15T07:36:05.394, value=val
          ue2
key4      column=cf1:c, timestamp=2025-05-15T07:36:08.459, value=val
          ue4
3 row(s)
Took 0.0288 seconds
hbase:011:0> put 'test', 'key1', 'cf1:a', 'value5'
Took 0.0078 seconds
hbase:012:0> scan 'test'
ROW COLUMN+CELL
key1      column=cf1:a, timestamp=2025-05-15T07:36:20.605, value=val
          ue5
key2      column=cf1:a, timestamp=2025-05-15T07:36:05.394, value=val
          ue2
key4      column=cf1:c, timestamp=2025-05-15T07:36:08.459, value=val
          ue4
3 row(s)
Took 0.0140 seconds
hbase:013:0>

```

```

hduser@ubuntu: ~/hbase_test/hbase-2.6.2/bin
21% 14 GB 75%
          ue1
key2      column=cf1:a, timestamp=2025-05-15T07:36:05.394, value=val
          ue2
key4      column=cf1:c, timestamp=2025-05-15T07:36:08.459, value=val
          ue4
3 row(s)
Took 0.0288 seconds
hbase:011:0> put 'test', 'key1', 'cf1:a', 'value5'
Took 0.0078 seconds
hbase:012:0> scan 'test'
ROW COLUMN+CELL
key1      column=cf1:a, timestamp=2025-05-15T07:36:20.605, value=val
          ue5
key2      column=cf1:a, timestamp=2025-05-15T07:36:05.394, value=val
          ue2
key4      column=cf1:c, timestamp=2025-05-15T07:36:08.459, value=val
          ue4
3 row(s)
Took 0.0140 seconds
hbase:013:0> disable 'test'
Took 0.3399 seconds
hbase:014:0> drop 'test'
Took 0.1331 seconds
hbase:015:0> exit
hduser@ubuntu:~/hbase_test/hbase-2.6.2/bin$

```

```

hduser@ubuntu: ~/hbase_test/hbase-2.6.2/bin
21% 14 GB 75%
3 row(s)
Took 0.0288 seconds
hbase:011:0> put 'test', 'key1', 'cf1:a', 'value5'
Took 0.0078 seconds
hbase:012:0> scan 'test'
ROW COLUMN+CELL
key1 column=cf1:a, timestamp=2025-05-15T07:36:20.605, value=va
ue5
key2 column=cf1:a, timestamp=2025-05-15T07:36:05.394, value=va
ue2
key4 column=cf1:c, timestamp=2025-05-15T07:36:08.459, value=va
ue4
3 row(s)
Took 0.0140 seconds
hbase:013:0> disable 'test'
Took 0.3399 seconds
hbase:014:0> drop 'test'
Took 0.1331 seconds
hbase:015:0> exit
hduser@ubuntu:~/hbase_test/hbase-2.6.2/bin$ sh stop-hbase.sh
stop-hbase.sh: 180: /home/hduser/hbase_test/hbase-2.6.2/bin/hbase-config.sh: Syn
tax error: "(" unexpected
hduser@ubuntu:~/hbase_test/hbase-2.6.2/bin$ ./stop-hbase.sh
stopping hbase.....
hduser@ubuntu:~/hbase_test/hbase-2.6.2/bin$

```

[ch 3-12]

```

hduser@ubuntu: ~/hbase_test/hbase-2.6.2/bin
22% 14 GB 74%
C 2025
Took 0.0012 seconds
hbase:001:0> create 'testdata', 'cf1', 'cf2'
Created table testdata
Took 1.0506 seconds
=> Hbase::Table - testdata
hbase:002:0> put 'testdata', 'rowA', 'cf1:A', 'val1'
Took 0.1207 seconds
hbase:003:0> put 'testdata', 'rowA', 'cf1:A', 'val1'
Took 0.0115 seconds
hbase:004:0> put 'testdata', 'rowC', 'cf2:C', 'val3'
Took 0.0044 seconds
hbase:005:0> put 'testdata', 'rowA', 'cf2:B', 'val4'
Took 0.0046 seconds
hbase:006:0> scan 'testdata'
ROW COLUMN+CELL
rowA column=cf1:A, timestamp=2025-05-15T07:39:19.643, value=val
1
rowA column=cf2:B, timestamp=2025-05-15T07:39:27.511, value=val
4
rowC column=cf2:C, timestamp=2025-05-15T07:39:23.710, value=val
3
2 row(s)
Took 0.0241 seconds
hbase:007:0>

```

```
hduser@ubuntu: ~/hbase_test/hbase-2.6.2/bin
28% 14 GB 74%
hbase:012:0> put 'testdata', 'rowA', 'cf1:A', 'val5'
Took 0.0076 seconds
hbase:013:0> scan 'testdata'
ROW COLUMN+CELL
rowA column=cf1:A, timestamp=2025-05-15T07:39:41.183, value=val5
rowA column=cf2:B, timestamp=2025-05-15T07:39:27.511, value=val4
rowC column=cf2:C, timestamp=2025-05-15T07:39:23.710, value=val3
2 row(s)
Took 0.0178 seconds
hbase:014:0> delete 'testdata', 'rowA', 'cf1:a'
Took 0.0163 seconds
hbase:015:0> scan 'testdata'
ROW COLUMN+CELL
rowA column=cf1:A, timestamp=2025-05-15T07:39:41.183, value=val5
rowA column=cf2:B, timestamp=2025-05-15T07:39:27.511, value=val4
rowC column=cf2:C, timestamp=2025-05-15T07:39:23.710, value=val3
2 row(s)
Took 0.0169 seconds
hbase:016:0>
```

```
hduser@ubuntu: ~/hbase_test/hbase-2.6.2/bin
26% 14 GB 74%
hbase:019:0>
hbase:020:0>
hbase:021:0> alter 'testdata', 'delete' => 'cf2'
Updating all regions with the new schema...
1/1 regions updated.
Done.
Took 1.7870 seconds
hbase:022:0> scan 'testdata'
ROW COLUMN+CELL
rowA column=cf1:A, timestamp=2025-05-15T07:39:41.183, value=val5
1 row(s)
Took 0.0160 seconds
hbase:023:0> alter 'testdata', 'cf2'
Updating all regions with the new schema...
1/1 regions updated.
Done.
Took 1.5975 seconds
hbase:024:0> scan 'testdata'
ROW COLUMN+CELL
rowA column=cf1:A, timestamp=2025-05-15T07:39:41.183, value=val5
1 row(s)
Took 0.0117 seconds
hbase:025:0>
```



```

hduser@ubuntu: ~/hbase_test/hbase-2.6.2/bin
23% 14 GB 74%
Done.
Took 1.5975 seconds
hbase:024:0> scan 'testdata'
ROW COLUMN+CELL
 rowA column=cf1:A, timestamp=2025-05-15T07:39:41.183, value=val
      5
1 row(s)
Took 0.0117 seconds
hbase:025:0> put 'testdata', 'rowA', 'cf1:A', 'val5'
Took 0.0140 seconds
hbase:026:0> put 'testdata', 'rowC', 'cf2:C', 'val3'
Took 0.0073 seconds
hbase:027:0> put 'testdata', 'rowA', 'cf2:B', 'val4'
Took 0.0056 seconds
hbase:028:0> scan 'testdata'
ROW COLUMN+CELL
 rowA column=cf1:A, timestamp=2025-05-15T07:40:22.784, value=val
      5
 rowA column=cf2:B, timestamp=2025-05-15T07:40:30.033, value=val
      4
 rowC column=cf2:C, timestamp=2025-05-15T07:40:27.417, value=val
      3
2 row(s)
Took 0.0172 seconds
hbase:029:0>

```

```

hduser@ubuntu: ~/hbase_test/hbase-2.6.2/bin
30% 14 GB 74%
ROW COLUMN+CELL
 rowA column=cf1:A, timestamp=2025-05-15T07:40:22.784, value=val
      5
 rowA column=cf2:B, timestamp=2025-05-15T07:40:30.033, value=val
      4
 rowC column=cf2:C, timestamp=2025-05-15T07:40:27.417, value=val
      3
2 row(s)
Took 0.0172 seconds
hbase:029:0>
hbase:030:0>
hbase:031:0> deleteall 'testdata', 'rowA'
Took 0.0035 seconds
hbase:032:0> scan 'testdata'
ROW COLUMN+CELL
 rowC column=cf2:C, timestamp=2025-05-15T07:40:27.417, value=val
      3
1 row(s)
Took 0.0092 seconds
hbase:033:0> truncate 'testdata'
Truncating 'testdata' table (it may take a while):
Disabling table...
Truncating table...
Took 1.0325 seconds
hbase:034:0>

```